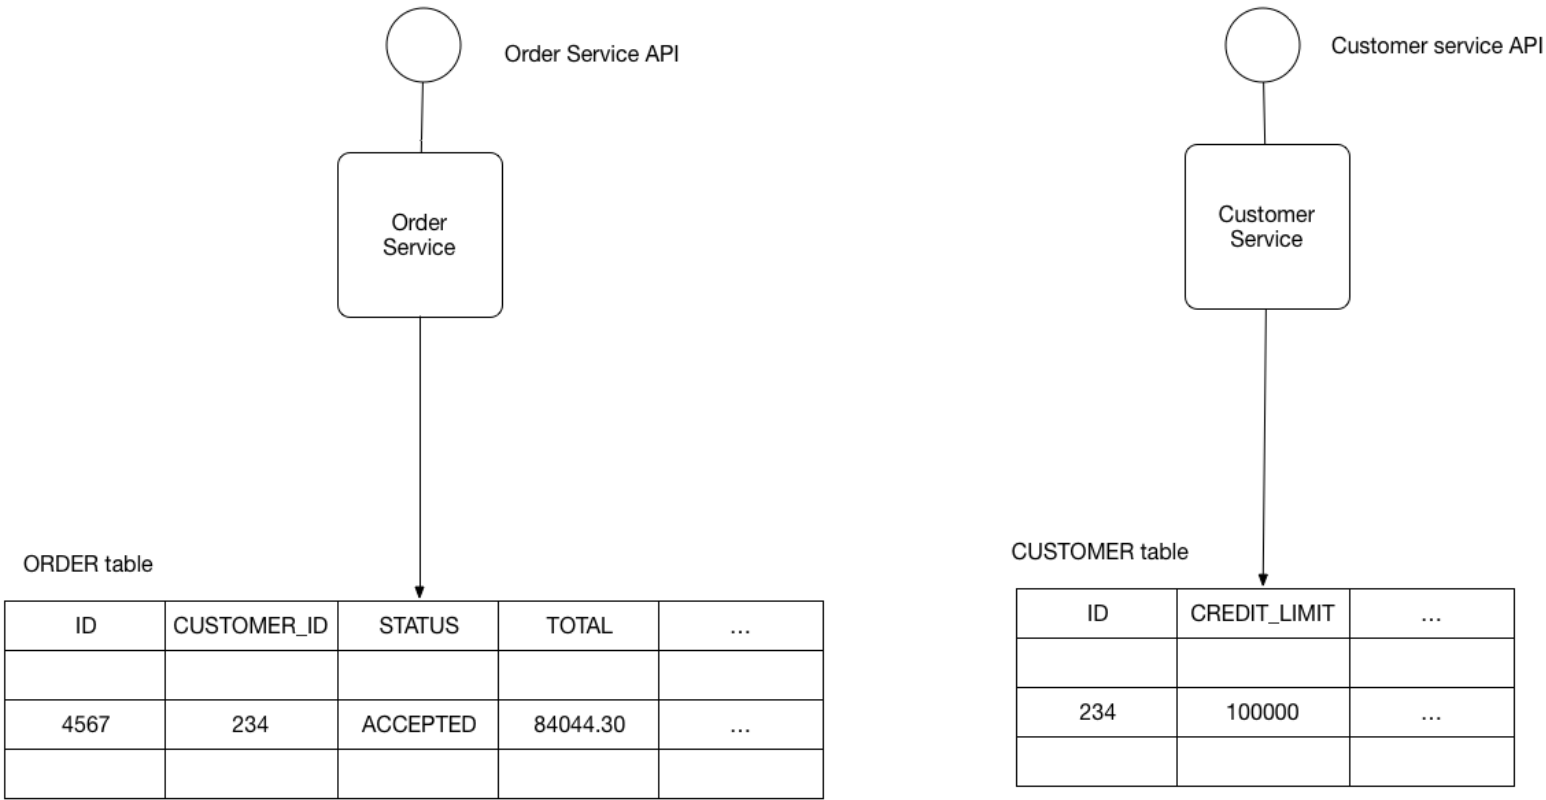


# Pattern: Shared database

## Context

Let's imagine you are developing an online store application using the Microservice architecture pattern (/patterns/microservices.html). Most services need to persist data in some kind of database. For example, the `Order Service` stores information about orders and the `Customer Service` stores information about customers.



## Problem

What's the database architecture in a microservices application?

## Forces

- Services must be loosely coupled so that they can be developed, deployed and scaled independently
- Some business transactions must enforce invariants that span multiple services. For example, the `Place Order` use case must verify that a new `Order` will not exceed the customer's credit limit. Other business transactions, must update data owned by multiple services.
- Some business transactions need to query data that is owned by multiple services. For example, the `View Available Credit` use must query the `Customer` to find the `creditLimit` and `Orders` to calculate the total amount of the open orders.
- Some queries must join data that is owned by multiple services. For example, finding customers in a particular region and their recent orders requires a join between customers and orders.
- Databases must sometimes be replicated and sharded in order to scale. See the `Scale Cube` (/articles/scalecube.html).

- Different services have different data storage requirements. For some services, a relational database is the best choice. Other services might need a NoSQL database such as MongoDB, which is good at storing complex, unstructured data, or Neo4J, which is designed to efficiently store and query graph data.

# Solution

Use a (single) database that is shared by multiple services. Each service freely accesses data owned by other services using local ACID transactions.

# Example

The `OrderService` and `CustomerService` freely access each other's tables. For example, the `OrderService` can use the following ACID transaction ensure that a new order will not violate the customer's credit limit.

```
BEGIN TRANSACTION
...
SELECT ORDER_TOTAL
  FROM ORDERS WHERE CUSTOMER_ID = ?
...
SELECT CREDIT_LIMIT
  FROM CUSTOMERS WHERE CUSTOMER_ID = ?
...
INSERT INTO ORDERS ...
...
COMMIT TRANSACTION
```

The database will guarantee that the credit limit will not be exceeded even when simultaneous transactions attempt to create orders for the same customer.

# Resulting context

The benefits of this pattern are:

- A developer uses familiar and straightforward ACID transactions to enforce data consistency
- A single database is simpler to operate

The drawbacks of this pattern are:

- Development time coupling - a developer working on, for example, the `OrderService` will need to coordinate schema changes with the developers of other services that access the same tables. This coupling and additional coordination will slow down development.
- Runtime coupling - because all services access the same database they can potentially interfere with one another. For example, if long running `CustomerService` transaction holds a lock on the `ORDER` table then the `OrderService` will be blocked.
- Single database might not satisfy the data storage and access requirements of all services.

# Related patterns

- Database per Service ([database-per-service.html](#)) is an alternative approach