

Pattern: Service Instance per VM

Context

You have applied the Microservice architecture pattern (/patterns/microservices.html) and architected your system as a set of services. Each service is deployed as a set of service instances for throughput and availability.

Problem

How are services packaged and deployed?

Forces

- Services are written using a variety of languages, frameworks, and framework versions
- Each service consists of multiple service instances for throughput and availability
- Service must be independently deployable and scalable
- Service instances need to be isolated from one another
- You need to be able to quickly build and deploy a service
- You need to be able to constrain the resources (CPU and memory) consumed by a service
- You need to monitor the behavior of each service instance
- You want deployment to be reliable
- You must deploy the application as cost-effectively as possible

Solution

Package the service as a virtual machine image and deploy each service instance as a separate VM

Examples

- Netflix packages each service as an EC2 AMI and deploys each service instance as a EC2 instance.

Resulting context

The benefits of this approach include:

- It's straightforward to scale the service by increasing the number of instances. Amazon Autoscaling Groups can even do this automatically based on load.
- The VM encapsulates the details of the technology used to build the service. All services are, for example, started and stopped in exactly the same way.
- Each service instance is isolated
- A VM imposes limits on the CPU and memory consumed by a service instance
- IaaS solutions such as AWS provide a mature and feature rich infrastructure for deploying and managing virtual machines. For example,
 - Elastic Load Balancer -
 - Autoscaling groups
 - ...

The drawbacks of this approach include:

- Building a VM image is slow and time consuming

Related patterns

- This pattern is a refinement of the Single Service per Host pattern ([single-service-per-host.html](#))
- The Service Instance per Container pattern ([service-per-container.html](#)) is an alternative solution
- The Serverless deployment pattern ([/patterns/deployment/serverless-deployment.html](#)) is an alternative solution.

Tweet

Follow [@MicroSvcArch](#)

Copyright © 2020 Chris Richardson • All rights reserved • Supported by Kong (<https://konghq.com/>).